

Investigating the Impact of REPA Loss on Training Speed and Generation Quality in Diffusion Models

Evgenii Pakhomov, Dmitrii Gromak, Sergei Kopylov, Miron Makhlin, Vladimir Novikov

HSE University, Methods of Optimization in Machine Learning course

Motivation

REPA accelerates diffusion training by aligning a transformer's hidden states with clean features from a frozen encoder (DINOv2). We ask *when* this helps, show that a constant alignment weight eventually hurts and compare effects on different single-domain datasets.

The Problem: Gradient Conflict

We track the cosine between the diffusion and REPA gradients, $c = \cos(\nabla_w \mathcal{L}_{\text{diff}}, \nabla_w \mathcal{L}_{\text{REPA}})$. It is *positive* at high noise but turns *negative* at low noise and late in training – there the auxiliary term pulls the update *against* the diffusion objective.




Figure 1. c vs. timestep t (high t = high noise). Note the negative region at low t .

REPA- Σ : Method & Guarantees

Discard *only* the part of the REPA gradient r anti-aligned with an EMA estimate $\bar{g}$ of the diffusion gradient g :

$$r^\Sigma = r - \frac{\langle r, \bar{g} \rangle}{\|\bar{g}\|_2^2} \bar{g}, \quad [x]_- = \min(x, 0).$$

Minimal change: projection of r onto $\{r' : \langle r', \bar{g} \rangle \geq 0\}$.

Non-interference: $\langle g + \lambda r^\Sigma, \bar{g} \rangle \geq \langle g, \bar{g} \rangle$ for all $\lambda \geq 0$ – the adapted gradient never works against the diffusion descent direction.

Toy Model: Whitened Linear Setting

Setting. Let $x_0 \in \mathbb{R}^d$ be a clean sample. We assume whitened features u with $\mathbb{E}[uu^\top] = I$ and a linear denoiser $\hat{x}_0 = Bu$. The diffusion loss $L_{\text{diff}}(B) = \frac{1}{2} \mathbb{E} \|Bu - x_0\|^2$ has the unique optimum $A = \mathbb{E}[x_0 u u^\top]$. We consider a teacher which projects onto an r -dim semantic subspace ($P^2 = P$, $Q = I - P$), giving $L_{\text{repa}}(B) = \frac{1}{2} \mathbb{E} \|Bu - Px_0\|^2$ – it keeps the semantic part Px_0 and *drops* the detail Qx_0 .

Result. Standard REPA minimises $L_{\text{diff}} + \lambda L_{\text{repa}}$ with optimum $B_\lambda^* = PA + \frac{1}{1+\lambda} QA$: the detail is shrunk, so it *cannot* reach A for $\lambda > 0$. REPA- Σ keeps A a stationary point for every λ – it provably reaches the optimal denoiser A while still using the alignment signal (under some extra assumptions).

Toy Experiments

Experiments supporting the Toy Model results above – empirically confirming that REPA- Σ reaches the optimal denoiser A while vanilla REPA stays biased.

Convergence across Three Datasets

This section compares the impact of REPA vs. REPA- Σ (w or w/o λ -anneal), against the diffusion baseline, on convergence speed (FID) across three single-domain datasets spanning a spatial-difficulty axis (the lower the dataset, the harder it is).

We plan to compare convergence curves between datasets.

On **CelebA**, surgery helps mid-training and λ -annealing late; combining both gives the best final FID.




Figure 2. CelebA – late-training FID. REPA- Σ + λ -anneal wins.

Feature Maps & Generation across Three Datasets

This section compares the impact of REPA vs. REPA- Σ using PCA feature maps and generations across the three datasets. In CelebA PCA maps show REPA preserves semantic, face-centred structure through the mid-network and under heavy noise, where the baseline collapses; we extend the comparison to all three domains.

We plan to compare feature maps between datasets.



Figure 3. CelebA – per-block PCA maps: image / Baseline / REPA.

LSUN Church
FID/KID convergence: baseline vs. REPA vs. REPA- Σ ($\pm \lambda$ -anneal), plus the gradient-compatibility curve.


LSUN Church
PCA feature maps & generation samples (baseline / REPA / REPA- Σ).

Comprehensive Cars
PCA feature maps & generation samples.

Comprehensive Cars (CompCars)
Same convergence comparison across all methods.

(1) Diagonal-matrix model with an independent-features teacher – closed-form convergence to A under REPA- Σ vs. biased REPA. (2) MLP denoiser on MNIST/CIFAR trained with a REPA-style teacher matching our setting.

Full Report



Full report – link to be added